

# Basics of Reinforcement Learning

Sina Baghal

---

## Abstract


This tutorial provides an introduction to the fundamentals of reinforcement learning. The main reference is the video lecture series by Sergey Levine.

## Contents

<b>1</b>	<b>What is RL?</b>	<b>1</b>
<b>2</b>	<b>Imitation Learning</b>	<b>3</b>
<b>3</b>	<b>REINFORCE</b>	<b>4</b>
<b>4</b>	<b>Variance Reduction</b>	<b>5</b>
<b>5</b>	<b>Bias Reduction</b>	<b>6</b>

## 1 What is RL?

**RL** In reinforcement learning, there is an *agent* and an *environment*. At time step  $t$ , the state is denoted by  $s_t$ . Given state  $s_t$ , the agent takes an action  $a_t$  resulting in a reward value  $r_t := r(s_t, a_t)$ .




**Policy** The agent's *policy* is parameterized by  $\pi_\theta$ , where  $\pi_\theta(\cdot \mid s_t)$  defines a probability distribution over possible actions at time  $t$ , given the state  $s_t$ .

**RL Goal** The goal of an RL algorithm is to maximize the *expected cumulative reward*:

$$\operatorname{argmax}_\theta \mathbb{E}_{\pi_\theta} \left[ \sum_{t=0}^T \gamma^t r(s_t, a_t) \right],$$

where  $0 \leq \gamma < 1$  and  $T$  are the discount factor and horizon resp. Notice that:

- More weight is placed on earlier steps.
- $\mathbb{E}_{\pi_\theta}$  is a smooth function of  $\theta$  where  $r$  itself may not be (e.g.,  $r \in \{\pm 1\}$ ).
- $s_t$  is independent of  $s_{t-1}$  (*Markov Property*).



**MDP** A *Markov Decision Process* (MDP) consists of a state space  $\mathcal{S}$  and an action space  $\mathcal{A}$ , along with a transition operator  $\mathcal{T}$  and a reward function  $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}_+$ . An MDP allows us to write a probability distribution over trajectories:

$$p_\theta(\tau) = p(s_1) \prod_{t=1}^T \pi_\theta(a_t|s_t) p(s_{t+1}|s_t, a_t), \quad \text{where } \tau = (s_1, a_1, \dots, s_T, a_T).$$

## 2 Imitation Learning

The analogous concept in reinforcement learning, compared to supervised learning, is called *imitation learning*, where the agent learns by mimicking expert actions. However, imitation learning often does not work well in practice due to the *distributional shift problem*. This arises because, in supervised learning, samples are assumed to be i.i.d., while in reinforcement learning the agent's past actions affect future states.

Assume that  $\pi^*$  is the expert policy and the learned policy  $\pi_\theta$  makes an error with probability at most  $\epsilon$  under the training distribution:

$$\Pr_{s_t \sim p_{\text{train}}} [\pi_\theta(s_t) \neq \pi^*(s_t)] \leq \epsilon.$$

Then,

$$p_\theta(s_t) = (1 - \epsilon)^t p_{\text{train}}(s_t) + (1 - (1 - \epsilon)^t) p_{\text{mistake}}(s_t).$$

Denote  $c_t(s_t, a_t) = 1_{\{a_t \neq \pi^*(s_t)\}} \in \{0, 1\}$ . Then the total number of times the policy  $\pi_\theta$  deviates from the optimal policy grows quadratically with  $T$ :

$$\begin{aligned} \mathbb{E}_{\pi_\theta} \left[ \sum_{t=0}^T c(s_t, a_t) \right] &= \sum_{t=0}^T \int p_\theta(s_t) c(s_t, a_t) ds_t \\ &= \sum_{t=0}^T (1 - \epsilon)^t \int p_{\text{train}}(s_t) c(s_t, a_t) ds_t + \sum_{t=0}^T (1 - (1 - \epsilon)^t) \int p_{\text{mistake}}(s_t) c(s_t, a_t) ds_t \\ &\leq \sum_{t=0}^T (1 - \epsilon)^t \epsilon + \sum_{t=0}^T 1 - (1 - \epsilon)^t \\ &\leq \sum_{t=0}^T (1 - \epsilon)^t \epsilon + 2\epsilon \sum_{t=0}^T t \\ &= \epsilon \cdot \mathcal{O}(T^2) \end{aligned}$$

This bound is achieved in the *tightrope walking* problem Figure 1, where the agent must learn to go straight; otherwise, it will enter unknown territory. Imitation learning can still be useful with some modifications, such as including bad actions along with corrective steps.



Figure 1: A tightrope walker.

### 3 REINFORCE

An MDP allows us to rewrite the goal of RL as the following optimization problem:

$$\operatorname{argmax}_{\theta} J(\theta) := \mathbb{E}_{\tau \sim p_{\theta}}[r(\tau)] = \int p_{\theta}(\tau) r(\tau) d\tau,$$

enabling a direct policy differentiation:

$$\begin{aligned} \nabla_{\theta} J(\theta) &= \int \nabla_{\theta} p_{\theta}(\tau) r(\tau) d\tau \\ &= \int p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau) r(\tau) d\tau \\ &= \mathbb{E}_{\tau \sim p_{\theta}} \nabla_{\theta} \log p_{\theta}(\tau) r(\tau) \\ &= \mathbb{E}_{\tau \sim p_{\theta}} \left( \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) \cdot \left( \sum_{t=1}^T r(s_t, a_t) \right) \quad \nabla_{\theta} p(s_{t+1} | s_t, a_t) = 0 \end{aligned}$$

We are now ready to state the first policy gradient method:

#### REINFORCE

1. Run the current policy  $N$  times to generate sample  $\tau_i$  for  $i = 1, \dots, N$ .
2. Compute the Monte Carlo estimate:

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left( \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \right) \cdot \left( \sum_{t=1}^T r(s_{i,t}, a_{i,t}) \right)$$

3. Apply Gradient Ascent:  $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$ .

## 4 Variance Reduction

One of the main issues with REINFORCE is the high variance in the reward term  $\sum_{t=1}^T r(s_{i,t}, a_{i,t})$ . In this section, we introduce some techniques to reduce this variance.

**Causality** As a first step toward variance reduction, we apply the *causality trick*:

Policy at time  $t'$  cannot impact reward at time  $t < t'$ .

Using which, the policy gradient is estimated as below:

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \left( \sum_{t'=t}^T r(s_{i,t'}, a_{i,t'}) \right)$$

The term  $\sum_{t'=t}^T r(s_{i,t'}, a_{i,t'})$  is referred to as the *reward-to-go*.

**Value Functions** The next idea is to replace the reward-to-go with a function estimator. To understand why this matters, see Figure 2. Notice two things: the ideal target for the reward-to-go function is the quantity  $Q(s_{i,t}, a_{i,t}) = \sum_{t'=t}^T \mathbb{E}_{\pi_{\theta}}[r(s_{i,t'}, a_{i,t'}) | s_{i,t}, a_{i,t}]$  rather than the single-sample estimate  $\sum_{t'=t}^T r(s_{i,t'}, a_{i,t'})$ . This represents the *value* of state  $s_{i,t}$  under the current policy where action  $a_{i,t}$  is taken at state  $s_{i,t}$ . Another advantage is that, as shown in Figure 2, if the state  $s'_{i,t}$  is quite close to  $s_{i,t}$  and  $p(s_{t+1} | s'_{i,t}, a'_{i,t}) \approx p(s_{t+1} | s_{i,t}, a_{i,t})$ , we expect their reward-to-go values to be similar. However, when working with a single-sample estimate, this relationship may easily be violated.




Figure 2: Value function fitting for variance reduction

**Baselines** Translation of the reward  $r \mapsto r - b$  can help reduce the variance. Assuming this translation,

$$\begin{aligned} \text{Var}[\nabla_{\theta} J(\theta)] &= \mathbb{E}_{\tau \sim p_{\theta}(\tau)} (\nabla_{\theta} \log p_{\theta}(\tau) (r(\tau) - b))^2 - (\mathbb{E}_{\tau \sim p_{\theta}(\tau)} \nabla_{\theta} \log p_{\theta}(\tau) (r(\tau) - b))^2 \\ &= \mathbb{E}_{\tau \sim p_{\theta}(\tau)} (\nabla_{\theta} \log p_{\theta}(\tau) (r(\tau) - b))^2 - (\mathbb{E}_{\tau \sim p_{\theta}(\tau)} \nabla_{\theta} \log p_{\theta}(\tau) r(\tau))^2 \end{aligned}$$

Table 1: Value Functions

Q-function (reward-to-go)	$Q^{\pi_\theta}(s_t, a_t)$	$\sum_{t'=t}^T \mathbb{E}_\theta[r(s_{t'}, a_{t'})   s_t, a_t]$
Value function	$V^{\pi_\theta}(s_t)$	$\mathbb{E}_{a_t \sim \pi_\theta(a_t   s_t)} [Q^{\pi_\theta}(s_t, a_t)]$
Advantage function	$A^{\pi_\theta}(s_t, a_t)$	$Q^{\pi_\theta}(s_t, a_t) - V^{\pi_\theta}(s_t)$

Thus, appropriate choice of  $b$  can reduce the variance. A proper choice is the expected value of  $Q$  function. Table 1 summarizes value functions used throughout. Note

$$\begin{aligned} Q(s_t, a_t) &= r(s_t, a_t) + \mathbb{E}_{s_{t+1} \sim p(\cdot | s_t, a_t)} V^{\pi_\theta}(s_{t+1}) \\ &\approx r(s_t, a_t) + V^{\pi_\theta}(s_{t+1}) \end{aligned}$$

Thus following policy gradient thus favors a lower variance.

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim p_\theta} \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot [r(s_t, a_t) + V^{\pi_\theta}(s_{t+1}) - V^{\pi_\theta}(s_t)]$$

**Discounts** The discount factor also helps reduce variance, as terms further in the horizon are weighted less. We then arrive at the following policy gradient:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim p_\theta} \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot [r(s_t, a_t) + \gamma \hat{V}_\phi^{\pi_\theta}(s_{t+1}) - \hat{V}_\phi^{\pi_\theta}(s_t)]$$

Here  $\hat{V}_\phi$  estimates  $V$ .

## 5 Bias Reduction

The policy gradient derived in the previous section, while enjoying low variance, is prone to higher bias. We tune this bias-variance trade-off as follows:  $n$ -step return estimator is:

$$\hat{A}_n^{\pi_\theta}(s_t, a_t) = \sum_{t'=t}^{t+n} \gamma^{t'-t} r(s_{t'}, a_{t'}) + \gamma^n \hat{V}_\phi(s_{t+n}) - \hat{V}_\phi(s_t)$$

For  $n = 1$ , we recover the previously mentioned policy gradient. As  $n \rightarrow +\infty$ , the bias is reduced while the variance increases. To manage this trade-off, we define

$$\begin{aligned} \hat{A}_{GAE}^{\pi_\theta} &= \sum_{n=1}^{+\infty} \lambda^{n-1} \hat{A}_n^{\pi_\theta} \\ &= \sum_{t'=t}^{+\infty} (\gamma \lambda)^{t'-t} \delta_{t'} \quad \delta_{t'} = r(s_{t'}, a_{t'}) + \gamma \hat{V}_\phi(s_{t'+1}) - \hat{V}_\phi(s_{t'}) \end{aligned}$$

We therefore arrive at the following policy gradient.

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p_{\theta}} \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot \hat{A}_{GAE}^{\pi_{\theta}}(s_t, a_t)$$